

Denoising Diffusion Probabilistic Model Implementation

02456 Deep Learning, Fall 2024, Final Project – Group 64

Eva Kaštelan (s232469), Jan Ljubas (s237214), Noa Margeta (s232470), Filip Penzar (s232452)

DDPM

Introduction

- Generative models
- Generate images by reversing a controlled noise process
- Two steps:
 - Forward diffusion process – transforming an image to noise*

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t \mathbf{I})$$

β_t is a variance schedule controlling the amount of noise added

- Reverse process – the model learns to undo noise addition in small steps*

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma_t^2 \mathbf{I})$$

The model predicts the noise $\epsilon_\theta(x_t, t)$ so we reparametrize $\mu_\theta(x_t, t)$:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(x_t, t) \right)$$

- In training the primary goal is to approximate the probability distribution of the data
- The loss function is then modelled as a variational lower bound
- Minimize the difference between the actual noise added at each step and the noise predicted by the model using MSE:

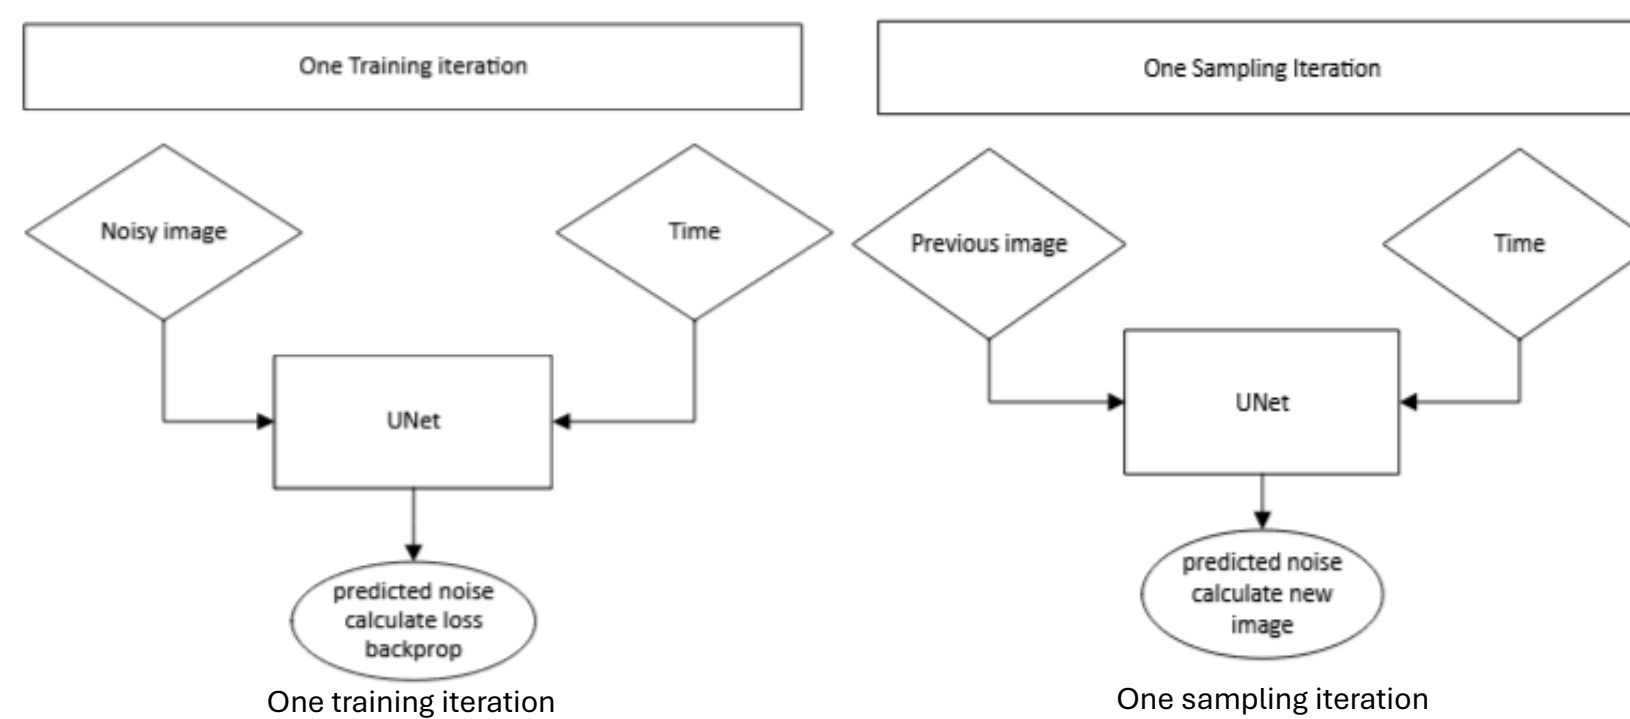
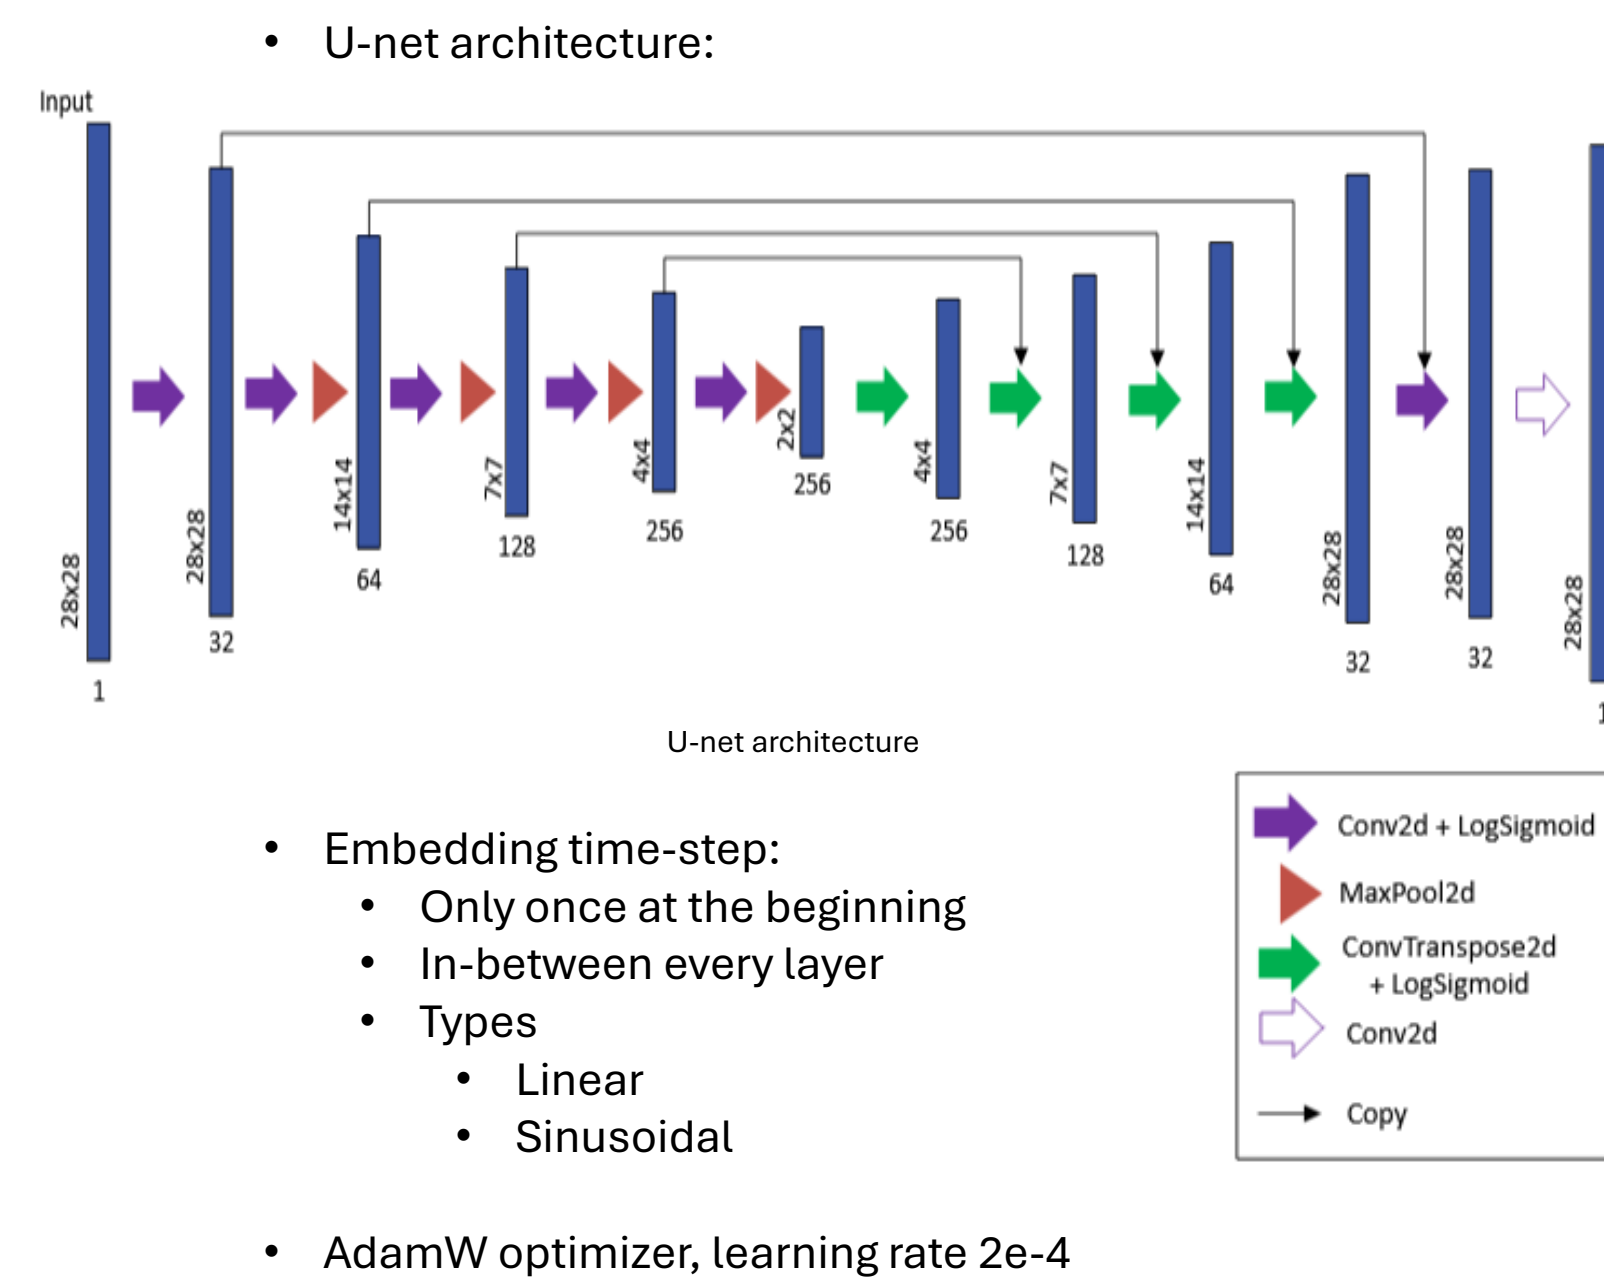
$$L = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t)\|^2]$$

- During sampling, the reverse process iteratively applies:

$$x_{t-1} = \mu_\theta(x_t, t) + \sigma_t z, \quad z \sim \mathcal{N}(0, \mathbf{I})$$


until a clean image is generated

Implementation




Results

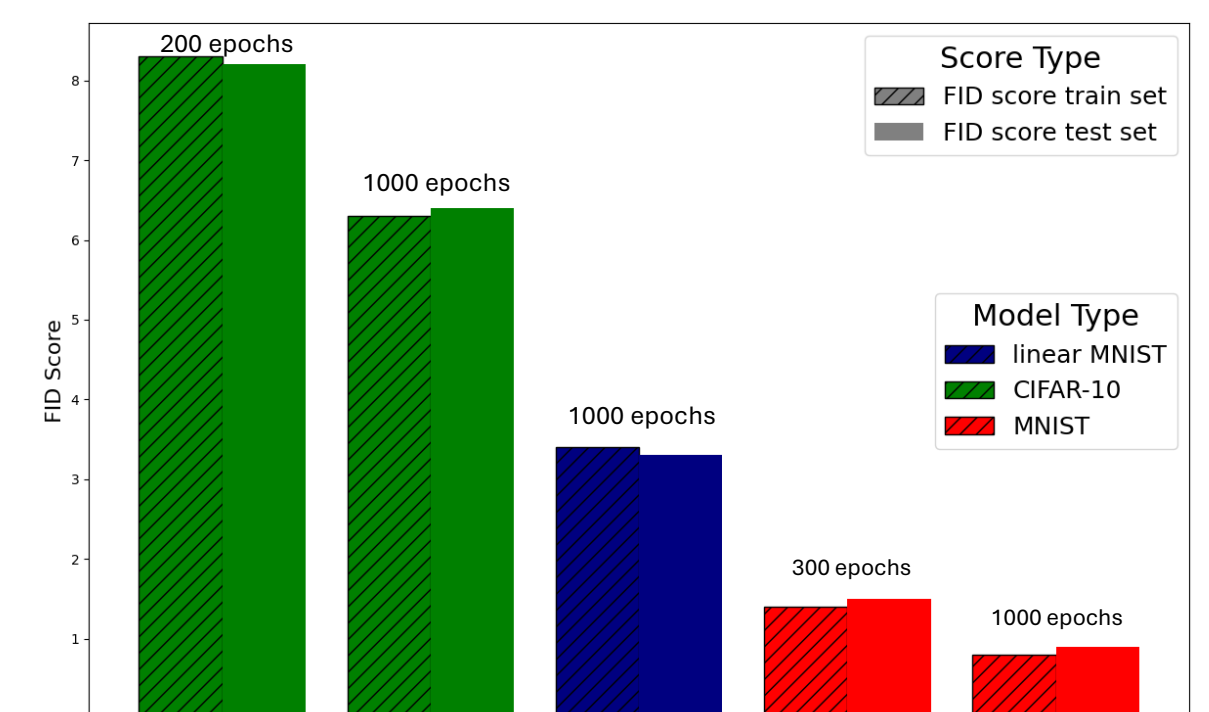
MNIST model samples



CIFAR-10 model samples



FID scores



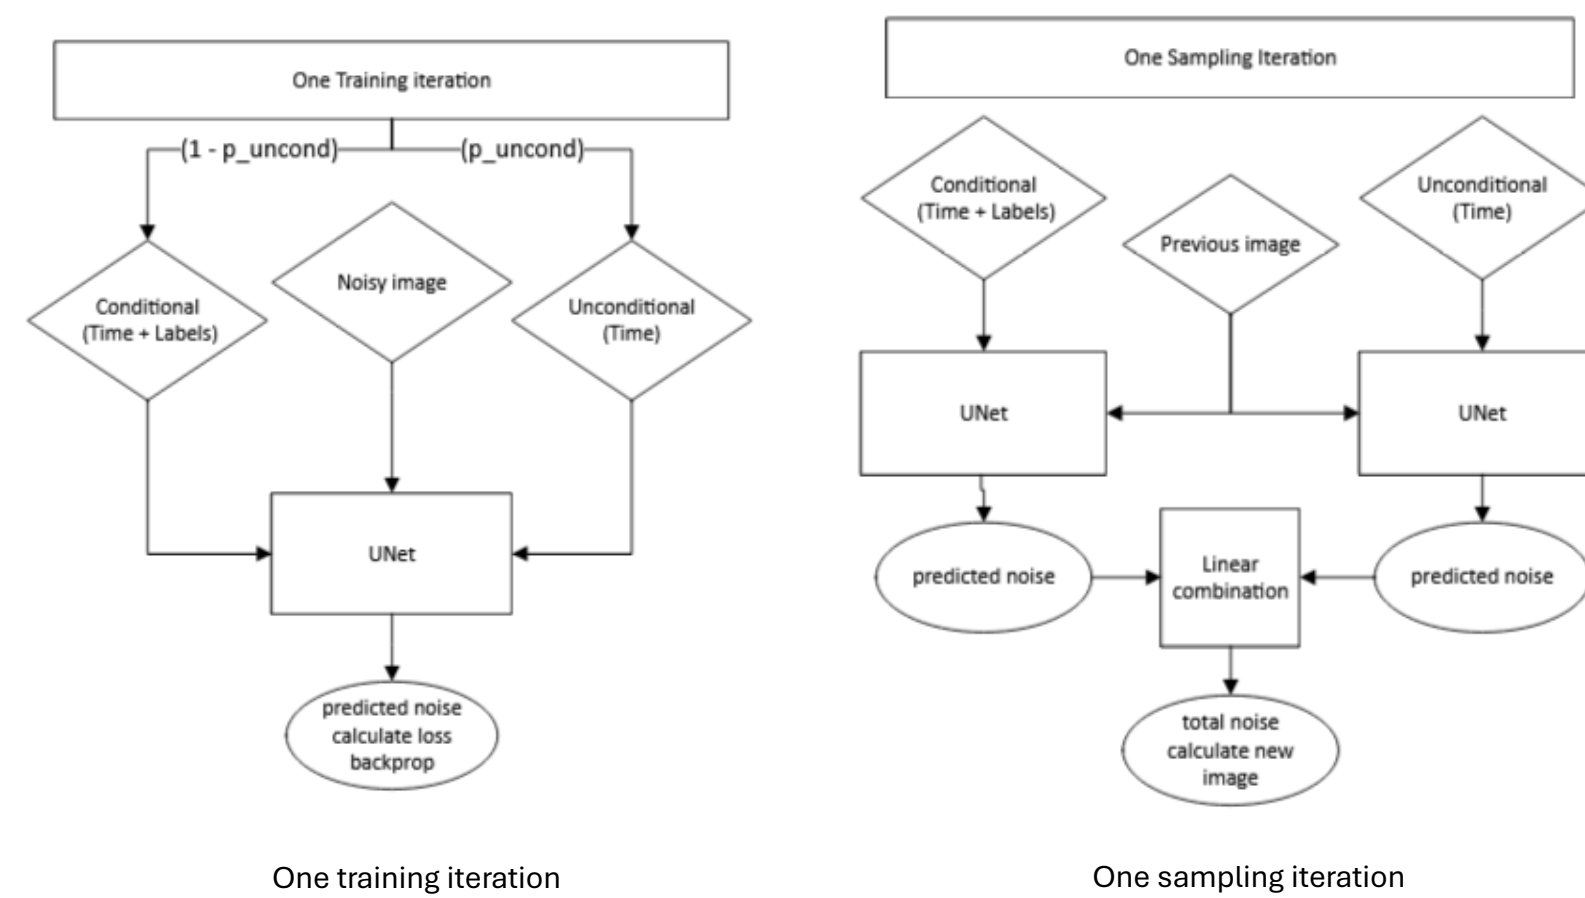
Conditional Diffusion Model

Introduction

- Goal of generating “low temperature” samples
- Classifier guidance
 - Uses an additional classifier
 - Mixes diffusion model’s score with classifier’s input gradient log probability
- Classifier-Free guidance
 - No need for an additional classifier
 - Joint training of a conditional and an unconditional model
 - Conditional model, $p(\mathbf{z}|\mathbf{c})$, with $\epsilon_\theta(\mathbf{z}_\lambda, \mathbf{c})$
 - Unconditional model, $p(\mathbf{z})$, with $\epsilon_\theta(\mathbf{z}_\lambda)$
 - Single neural network for both, unconditional = conditional with $\mathbf{c} = \emptyset$
- Training
 - Training as a conditional model with p_{uncond} probability of training the unconditional model
- Sampling
 - Linear combination of the conditional and the unconditional models with guidance strength (w)
 - $\tilde{\epsilon}_\theta(\mathbf{z}_\lambda, \mathbf{c}) = (1 + w)\epsilon_\theta(\mathbf{z}_\lambda, \mathbf{c}) - w\epsilon_\theta(\mathbf{z}_\lambda)$


Implementation

- Same architecture as our DDPM
- Sinusoidal time embedding
- Embedding for labels (28 dimensions)
- Time embedding and label embedding are added together, and appended to the model in forward process (conditional)
- Only time embedding is used for unconditional
- $p_{uncond} = 0.1$
- $w = 3$

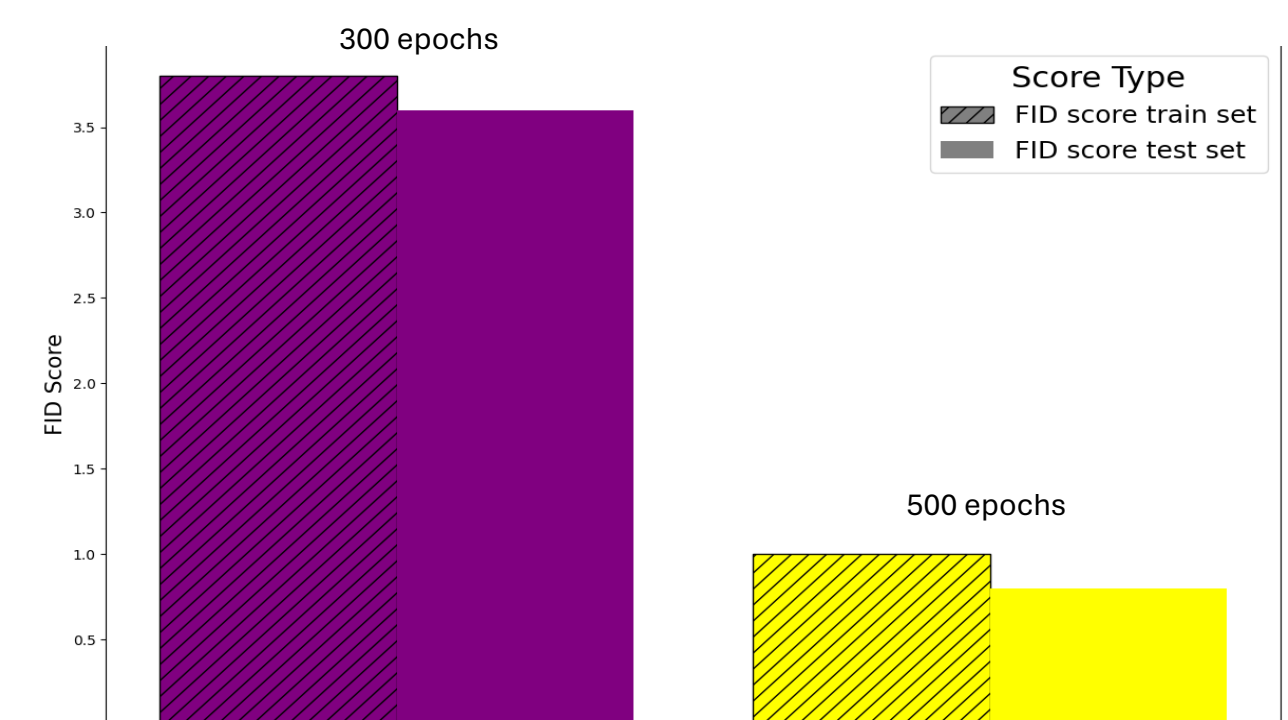


Results


CFG-MNIST model samples




FID scores



Unconditional MNIST progressive generation:



Conditional MNIST progressive generation (label=4):



References

- Jonathan Ho, Ajay Jain, and Pieter Abbeel. "Denoising Diffusion Probabilistic Models." *Proceedings of the 34th Conference on Neural Information Processing Systems (NeurIPS 2020)*
- Jonathan Ho and Tim Salimans. "Classifier-Free Diffusion Guidance." *NeurIPS 2021 Workshop on Deep Generative Models and Downstream Applications*
- Paszke, A., Gross, S., Chintala, S., & Chanan, G. (2016). "PyTorch: An imperative style, high-performance deep learning library"
- Outlier. (2022, June 6). Diffusion models | Paper explanation | Math explained [Video]. YouTube. <https://www.youtube.com/watch?v=HoKDTa5jHvg>